

Data controls in the OpenAI platform

Understand how OpenAI uses your data, and how you can control it.

 Copy Page

Understand how OpenAI uses your data, and how you can control it.

Your data is your data. As of March 1, 2023, data sent to the OpenAI API is not used to train or improve OpenAI models (unless you explicitly opt in to share data with us).

Types of data stored with the OpenAI API

When using the OpenAI API, data may be stored as:

- **Abuse monitoring logs:** Logs generated from your use of the platform, necessary for OpenAI to enforce our [Usage Policies](#) and agreements and mitigate harmful uses of AI.
- **Application state:** Data persisted from some API features in order to fulfill the task or request.

Data retention controls for abuse monitoring

Abuse monitoring logs may contain certain customer content, such as prompts and responses, as well as metadata derived from that customer content, such as classifier outputs. By default, abuse monitoring logs are generated for all API feature usage and retained for up to 30 days, unless longer retention is required by law, or is reasonably necessary to protect our services or any third party from harm.

Eligible customers may have their customer content excluded from these abuse monitoring logs, subject to the limitations below, by getting approved for the [Zero Data Retention](#) or [Modified Abuse Monitoring](#) controls. Currently, these controls are subject to prior approval by OpenAI and acceptance of additional requirements.

Approved customers may select between Modified Abuse Monitoring or Zero Data Retention for their API Organization or project.

Ask AI

Loading docs agent...

Eyes Off

For customers approved for Zero Data Retention or Modified Abuse Monitoring, we reserve the right to make models ineligible for Zero Data Retention or Modified Abuse Monitoring for specific customers, as notified in advance to the impacted customers in writing. In this instance, customer content will be retained in abuse monitoring logs, but such content will be excluded from human review unless required by applicable law. For customers who have executed an OpenAI Business Associate and Healthcare Addendum, once your org ID is provisioned with Eyes Off, BAA-eligible endpoints can be used for processing PHI, even if data is retained.

Safety Retention

For customers approved for Zero Data Retention or Modified Abuse Monitoring, we reserve the right to make models ineligible for Zero Data Retention or Modified Abuse

Docs agent

Loading docs agent...

OpenAI may retain application state when using certain OpenAI endpoints.

Retention enabled.

Endpoint	Data used for training	Abuse monitoring retention	Application state retention	Zero Data Retention eligible	Eyes Off and Safety Retention eligible
/v1/chat/completions	No	30 days	None, see below for exceptions	Yes, see below for limitations	Yes, see below for limitations
/v1/responses	No	30 days	None, see below for exceptions	Yes, see below for limitations	Yes, see below for limitations

Docs agent

Loading docs agent...

/v1/vector_stores	No	30 days	Until deleted	No	No
/v1/images/generations	No	30 days	None	Yes, see below for limitations	No
/v1/images/edits	No	30 days	None	Yes, see below for limitations	No
/v1/embeddings	No	30 days	None	Yes	No
/v1/audio/transcriptions	No	None	None	Yes	No
/v1/audio/translations	No	None	None	Yes	No

Docs agent

Loading docs agent...

7.17. End of completions

- Audio outputs application state is stored for 1 hour to enable [multi-turn conversations](#).
- When Zero Data Retention is enabled for an organization, the `store` parameter will always be treated as `false`, even if the request attempts to set the value to `true`.
- See [image and file inputs](#).
- Prompt caching may store encrypted key/value tensors in GPU-local storage as application state. This data is stored on the local GPU machines and is not retained after the 24-hour expiration. For `gpt-5.5` and `gpt-5.5-pro`, setting `prompt_cache_retention` to `in_memory` returns an error. For GPT-5.6 models and later model families, `prompt_cache_options.ttl` controls the minimum cache lifetime, not this maximum application-state retention period. To learn more, see the [prompt caching guide](#).

Docs agent

Loading docs agent...

application state. This data is stored on the local GPU machines and is not retained after the 24-hour expiration. For `gpt-5.5` and `gpt-5.5-pro`, setting `prompt_cache_retention` to `in_memory` returns an error. For GPT-5.6 models and later model families, `prompt_cache_options.ttl` controls the minimum cache lifetime, not this maximum application-state retention period. To learn more, see the [prompt caching guide](#).

- When Zero Data Retention is not enabled for an organization, all queries use extended prompt caching for all supported models.
- For server-side compaction, no data is retained when `store="false"`.
- We support **Skills** in two form factors, both local execution and hosted container-based execution. Hosted skills follow the same container lifecycle as hosted shell: mounted skills and container files remain available while the container is active and are discarded when the container expires or is deleted.

Docs agent

Loading docs agent...

None as described in [Configuring data retention controls](#) to use `/v1/videos` with that project.

Image and file inputs

Images and files may be uploaded as inputs to `/v1/responses` (including when using the Computer Use tool), `/v1/chat/completions`, and `/v1/images`. Image and file inputs are scanned for CSAM content upon submission. If the classifier detects potential CSAM content, the image will be retained for manual review, even if Zero Data Retention, Modified Abuse Monitoring, or Eyes Off is enabled.

Web Search

Web Search with live internet access is not HIPAA eligible and is not covered by a BAA. Web Search in offline/cache-only mode (`external_web_access: false`) is

Docs agent

Loading docs agent...

Data residency does not apply to system data, which may be processed and stored outside the selected region. System data means account data, metadata, and usage data that do not contain Customer Content, which are collected by the services and used to manage and operate the services, such as account information or profiles of end users that directly access the services (for example, your personnel), analytics, usage statistics, billing information, support requests, and structured output schema.

Sub-processors and regional request processing

OpenAI uses sub-processors to provide its services. For requests sent to `us.api.openai.com` or `eu.api.openai.com`, OpenAI uses Cloudflare Regional Services so that TLS termination and HTTPS decryption occur within the selected processing region.

Docs agent

Loading docs agent...

For requests to projects with data residency configured, add the domain prefix as defined in the table below to each request.

Select a processing region per request

As an alternative to creating a region-specific project, you can select regional processing for an individual request by using the prefixed domain with an API key from a project having Global geography.

Existing eligibility and data retention control requirements still apply. The selected endpoint and model must also support regional processing, as shown in the table below.

The following example reuses one client and an API key from a Global project for global, US, and EU requests:

Docs agent

Loading docs agent...

The following models and API services are eligible for data residency today for the regions specified below.

Use **Support by region** to compare regional capabilities and expand the services available in each region. Use **API Endpoint, tool and model support** for complete model lists and a detailed service view. Support for regional storage does not imply support for regional processing.

Support by region

Compare regional capabilities and services.

All regions

United States

us.api.openai.com

Storage

Processing

Docs agent

Loading docs agent...

/v1/chat/completions Storage Processing

Show 21 more services

Australia Storage

au.api.openai.com

Storage	Processing
Yes	No

Requires MAM or ZDR

Modes: Text, Audio, Voice, Image*

SUPPORTED SERVICES

/v1/audio/transcriptions, /v1/audio/translations, /v1/audio/speech Storage

Docs agent

Loading docs agent...

Storage

Yes

Processing

No

Requires MAM or ZDR

Modes: Text, Audio, Voice, Image*

SUPPORTED SERVICES

/v1/audio/transcriptions, /v1/audio/translations, /v1/audio/speech Storage

/v1/batches Storage

/v1/chat/completions Storage

Show 17 more services

India

Storage

Docs agent

Loading docs agent...

/v1/audio/transcriptions, /v1/audio/translations, /v1/audio/speech

Storage

/v1/batches

Storage

/v1/chat/completions

Storage

Show 17 more services

South Korea

Storage

kr.api.openai.com

Storage

Processing

Yes

No

Requires MAM or ZDR

Modes: Text, Audio, Voice, Image*

Docs agent

Loading docs agent...

United Arab Emirates

Storage Processing

ae.api.openai.com

Storage Processing
Yes Yes

Requires MAM or ZDR

Modes: Text, Audio, Voice, Image*

SUPPORTED SERVICES

/v1/audio/transcriptions, /v1/audio/translations, /v1/audio/speech Storage

/v1/batches Storage

/v1/chat/completions Storage Processing

Docs agent

Loading docs agent...

/v1/batches	Batches	gpt-5.5-pro-2026-04-23 gpt-5.4-pro-2026-03-05 gpt-5.2-pro-2025-12-11 ► + 28 model snapshots	—
/v1/chat/completions	Chat Completions	gpt-5.6-sol gpt-5.6-terra gpt-5.6-luna ► + 22 model snapshots	—

Docs agent

Loading docs agent...

		gpt-image-1 gpt-image-1.5 ► + 1 model snapshot	
/v1/moderations	Moderation	omni-moderation-latest	—
/v1/realtime	Realtime	gpt-realtime gpt-realtime-1.5 gpt-realtime-mini ► + 3 model snapshots	—

Docs agent

Loading docs agent...

File Search	Tools	Supported	—
File Uploads	Files	Supported	Supports file uploads
Remote MCP server tool	Tools	Supported	Must be configured to support remote tools
Scale Tier	Other	Supported	—

Docs agent

Loading docs agent...

/v1/realtime

Tracing is not currently EU data residency compliant for `/v1/realtime`.

Enterprise Key Management (EKM)

Enterprise Key Management (EKM) allows you to encrypt your customer content at OpenAI using keys managed by your own external Key Management System (KMS).

Once configured, EKM applies to any [application state](#) created during your use of the platform. See the [EKM help center article](#) for more information about how EKM works, and how to integrate with your KMS provider.

EKM limitations

Docs agent

Loading docs agent...